

# Data Base Systems

## VU 184.686, WS 2020

### Transaction Management

Anela Lolić

Institute of Logic and Computation, TU Wien



FAKULTÄT  
FÜR INFORMATIK

Faculty of Informatics

# Acknowledgements

The slides are based on the slides (in German) of [Sebastian Skritek](#).

The content is based on [Chapter 9](#) of  
(Kemper, Eickler: Datenbanksysteme – Eine Einführung).  
Many examples and illustrations are taken from there.

For related literature in English see [Chapter 16](#) of  
(Ramakrishnan, Gehrke: Database Management Systems).

# Transaction Management

## 1. Architecture of a DBMS

## 2. Transactions

### 2.1 Definition

### 2.2 Requirements and Features

## 3. Transaction Management in SQL

# Overview

## 1. Architecture of a DBMS


## 2. Transactions

### 2.1 Definition

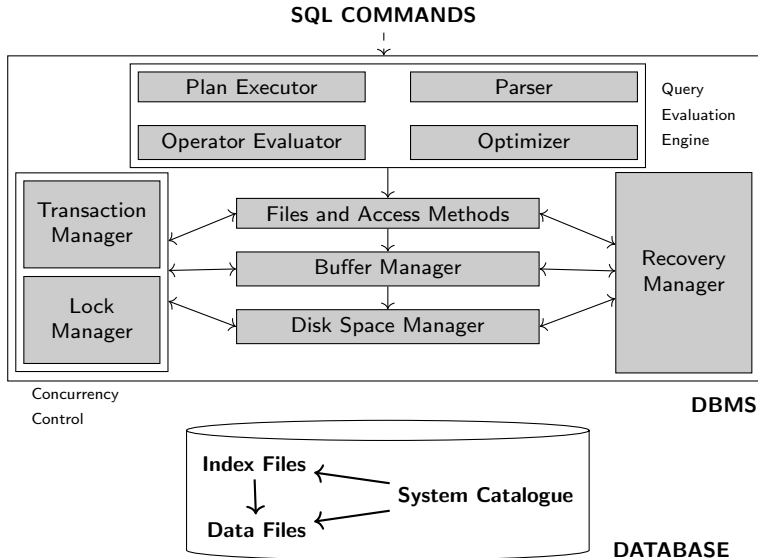
### 2.2 Requirements and Features

## 3. Transaction Management in SQL

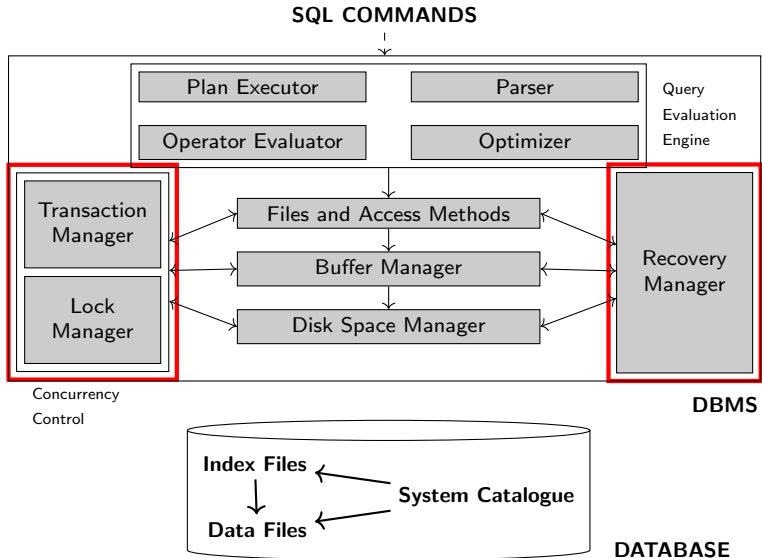
# Architecture of a DBMS



# Architecture of a DBMS

**DATABASE**

# Architecture of a DBMS

**DATABASE**

# Concurrency Control

## **transaction manager:**

- controls the processing of transactions

## **locking manager:**

- maintains lock requests on data objects (tuple, page, ...)
- fulfils lock requests for data objects as soon as they are available



# Recovery Manager

during operation:

- management of log-file

during recovery after system failure or crash:

- recovery of a consistent state, i.e.
  - redo of all lost operations of successfully completed transactions
  - undo of all operations of not successfully completed transactions

# Overview

## 1. Architecture of a DBMS

## 2. Transactions

### 2.1 Definition

### 2.2 Requirements and Features

## 3. Transaction Management in SQL

# Overview

## 1. Architecture of a DBMS

## 2. Transactions

### 2.1 Definition

### 2.2 Requirements and Features

## 3. Transaction Management in SQL

# What is a Transaction?

## Example

a typical transaction in a bank:

account  $A$ , account  $B$

- |  |                         |
|--|-------------------------|
| 1. read the balance from $A$ into the variable $a$ : | <b>read</b> ( $A,a$ );  |
| 2. reduce the balance by 50 €:                       | $a := a - 50$ ;         |
| 3. write the new balance into the database:          | <b>write</b> ( $A,a$ ); |
| 4. read the balance of $B$ into the variable $b$ :   | <b>read</b> ( $B,b$ );  |
| 5. increase the balance by 50 €:                     | $b := b + 50$ ;         |
| 6. write the new balance into the database:          | <b>write</b> ( $B,b$ ); |

# Transaction

## Definition (transaction)

A **transaction** is a **sequence of database operations** that transform the database from a consistent state into a consistent state. The execution of a transaction is **atomic** (= (logical) not interruptible) i.e.

- **as unit** sound and
- **without interference** by other transactions.

# Operations of a Transaction

for **data handling**:

$\text{read}(A, a)$  reads the value of a field  $A$  into a *local variable*  $a$

$\text{write}(A, a)$  writes the value  $a$  into the field  $A$

# Operations of a Transaction

for **data handling**:

`read( $A, a$ )` reads the value of a field  $A$  into a *local variable*  $a$

`write( $A, a$ )` writes the value  $a$  into the field  $A$

for **transaction control**:

`BOT` (begin of transaction)

marks the beginning of a transaction

`commit` initiates the successful termination of a transaction

`abort` initiates the abortion of a transaction,  
causes the DBMS to reset the database to the state  
it was in before the transaction was executed

# Operations of a Transaction

**additional** operations for transaction control:

**define savepoint** defines a savepoint to which the (still active) transaction can be reset  
complete abortion via **abort** is still possible

**roll-back to save point** initiates the reset of the active transaction to a savepoint  
based on the DBMS reset to the last savepoint or to older savepoints possible



# Schedule

**schedule** describes the order of elementary operations during an interleaved execution of several transactions

## Example

	$T_1$	$T_2$	$T_3$
1	BOT		
2		BOT	
3			BOT
4		read( $B, b_1$ )	
5	write( $A, a_1$ )		
6			read( $C, c_1$ )
7	commit		
8			abort

# Termination of a Transaction

- 1 successful termination through a **commit**
- 2 unsuccessful termination (requires subsequent reset)
  - initiated by user via **abort** (or **roll-back**)
  - initiated by DBMS based on an **error**

# Termination of a Transaction

- 1 successful termination through a **commit**
- 2 unsuccessful termination (requires subsequent reset)
  - initiated by user via **abort** (or **roll-back**)
  - initiated by DBMS based on an **error**

**BOT** $op_1$  $op_2$  $\vdots$  $op_n$ **commit****BOT** $op_1$  $op_2$  $\vdots$  $op_m$ **abort****BOT** $op_1$  $op_2$  $\vdots$  $op_k$  **error**

# Overview

## 1. Architecture of a DBMS

## 2. Transactions

### 2.1 Definition

### 2.2 Requirements and Features

## 3. Transaction Management in SQL

# Requirements of Transactions– ACID

**Atomicity** a transaction is the smallest, not further composable, unit (“everything or nothing at all”).

**Consistency** a transaction after completion leaves the database in a consistent state

**Isolation** concurrently executed transactions shall not influence each other


**Durability** the effects of a successfully completed transactions are not lost  
(even when system errors occur)

# Components of Transaction Management


**concurrency control:** ensure isolation; controls concurrency  
concurrency is necessary for performance

**recovery:** ensures atomicity and durability  
guarantees „everything or nothing at all“ and makes  
sure that modifications of a successfully terminated  
transactions are not lost even when system errors  
occur

# Atomicity and Durability




# Atomicity and Durability



- transaction  $T_1$  has to exist after recovery



# Atomicity and Durability



- transaction  $T_1$  has to exist after recovery
- transaction  $T_2$ : all modifications to the database through  $T_2$  have to be removed after recovery

# Learning Objectives

- What are transactions?
- Are they important?
- Which operations within a transaction are there?
- What are the possible ends of transactions?
- What do the ACID-features tell us and what is their effect to the transaction management?

# Overview

## 1. Architecture of a DBMS

## 2. Transactions

### 2.1 Definition

### 2.2 Requirements and Features

## 3. Transaction Management in SQL

# Transaction Management in SQL

*remark:* differences in details between the different DBMS

## start of a transaction:

- implicitly through an instruction
- explicitly through SQL commands

```
START TRANSACTION  
BEGIN [ WORK | TRANSACTION ]
```

# Transaction Management in SQL

## termination of a transaction:

- implicitly (“auto commit”)
- explicitly

```
COMMIT [ WORK | TRANSACTION ]
```

- as long as no problems occur (e.g. consistency violations) the modifications are performed

```
ROLLBACK [ WORK | TRANSACTION ]
```

```
ABORT [ WORK | TRANSACTION ]
```

- modifications are reset
- DBMS has to guarantee the successful execution

# Implicit Transaction End

## **implicit** `commit`

- at implicit beginning and `AUTO COMMIT = ON`:  
after each DML/DDI command
- based on DMBS always after  
DDL (`CREATE TABLE, ...`) and DCL (`GRANT, ...`)  
commands

## **implicit** `roll-back`

- at problems like system crash, disconnections, consistency  
violations at `commit, ...`

**good stile:** whenever possible, end transactions **explicitly**

# Transaction Management in SQL – Savepoints

## ■ **SAVEPOINT** *sp\_name*

- defines reset point within running transaction
- enables reset of transaction to this point
- “economical” use is recommended:  
Can a big transaction be split in smaller ones?

## ■ **ROLLBACK** [WORK | **TRANSACTION**] TO [SAVEPOINT] *sp\_name*

- resets modifications of transactions since *sp\_name*

## ■ **RELEASE** [SAVEPOINT] *sp\_name*

- deletes save point (no other effect)