

# Data Base Systems

VU 184.686, WS 2021

Recovery

Institute of Logic and Computation, TU Wien



FAKULTÄT  
FÜR INFORMATIK

Faculty of Informatics

# Acknowledgements

The slides are based on the slides (in German) of [Sebastian Skritek](#).

The content is based on [Chapter 10](#) of  
(Kemper, Eickler: Datenbanksysteme – Eine Einführung).  
Many examples and illustrations are taken from there.

For related literature in English see [Chapter 18](#) of  
(Ramakrishnan, Gehrke: Database Management Systems).

# Recovery (Error Handling and Logging)

1. Memory Management
2. Error Categories
3. Taken System Configuration
4. Logging
5. Recovery after an Error
6. Checkpoints

# Overview

## 1. Memory Management

### 1.1 Storage Hierarchy and Memory Management

### 1.2 Management of the DBMS-Buffer

### 1.3 Introducing Strategy

## 2. Error Categories

## 3. Taken System Configuration

## 4. Logging

## 5. Recovery after an Error

## 6. Checkpoints

# Overview

## 1. Memory Management

### 1.1 Storage Hierarchy and Memory Management

### 1.2 Management of the DBMS-Buffer

### 1.3 Introducing Strategy

## 2. Error Categories

## 3. Taken System Configuration

## 4. Logging

## 5. Recovery after an Error

## 6. Checkpoints

# Storage Hierarchy

available storage media typically form a **hierarchy**

■ **trade-off:**

- fast and expensive vs. slow and cheap
- fast and volatile vs. slow and persistent
- fast and small vs. slow and big

■ **“access gap”** between the stages

# Storage Hierarchy

available storage media typically form a **hierarchy**

- **trade-off:**

- fast and expensive vs. slow and cheap
- fast and volatile vs. slow and persistent
- fast and small vs. slow and big

- **“access gap”** between the stages

(simplified) assumption in the course: **Two-Stage Storage Hierarchy**

- **database buffer** (fast access, volatile)
- **background memory** (slow access, persistent)

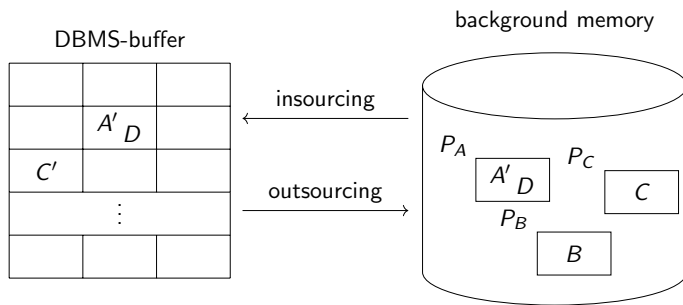
# Memory Organization: Pages

- memory split in **pages**
- dataset are represented by pages
  - simplified assumption in course: data  $A$  on page  $P_A$
- **granularity level** in which the data is moved between buffer and background memory



# Memory Organization: Pages

- memory split in **pages**
- dataset are represented by pages
  - simplified assumption in course: data  $A$  on page  $P_A$
- **granularity level** in which the data is moved between buffer and background memory



# Overview

## 1. Memory Management

### 1.1 Storage Hierarchy and Memory Management

### 1.2 Management of the DBMS-Buffer

### 1.3 Introducing Strategy

## 2. Error Categories

## 3. Taken System Configuration

## 4. Logging

## 5. Recovery after an Error

## 6. Checkpoints

# Replacement of Buffer Pages

- no outsourcing during access (FIX)
- after that based on chosen strategy
- *dirty* pages: modified pages in buffer

# Replacement of Buffer Pages

- no outsourcing during access (FIX)
- after that based on chosen strategy
- *dirty* pages: modified pages in buffer

**strategy** for the replacement of buffer pages:

*steal* every non-fixed page is in principle a candidate for outsourcing

¬*steal* pages that are modified by a still active transaction may not be outsourced

# Outsourcing of Buffer Pages

**strategy** for the outsourcing of modified buffer pages after a successful transaction:

- force* modified pages are outsourced at the end of the transaction
- ¬*force* outsourcing at the end of a transaction is not forced, *dirty pages may stay in buffer*

# Overview

## 1. Memory Management

### 1.1 Storage Hierarchy and Memory Management

### 1.2 Management of the DBMS-Buffer

### 1.3 Introducing Strategy

## 2. Error Categories

## 3. Taken System Configuration

## 4. Logging

## 5. Recovery after an Error

## 6. Checkpoints

# Introducing Strategy

**update-in-place** (“direct” introducing strategy):

- every page has **exactly one** “place” in the background memory
- a page is copied at this place in the outsourcing process
- the old state of the page is overwritten

# Introducing Strategy

## **update-in-place** (“direct” introducing strategy):

- every page has **exactly one “place”** in the background memory
- a page is copied at this place in the outsourcing process
- the old state of the page is overwritten

## **twin-block-method** (“indirect” introducing strategy):

- for every page **two copies** in the background memory + **bit** indicating **current** copy
- page is copied at **current copy** during outsourcing



# Introducing Strategy

## **update-in-place** (“direct” introducing strategy):

- every page has **exactly one “place”** in the background memory
- a page is copied at this place in the outsourcing process
- the old state of the page is overwritten

## **twin-block-method** (“indirect” introducing strategy):

- for every page **two copies** in the background memory + **bit** indicating **current** copy
- page is copied at **current copy** during outsourcing

## **shadow memory concept** (“indirect” introducing strategy):

- there are two copies only for **modified pages**

# Overview

1. Memory Management
2. Error Categories
3. Taken System Configuration
4. Logging
5. Recovery after an Error
6. Checkpoints

# Three Error Categories

## **local error** (in a transaction)

- only one transaction is affected
- modifications caused by this transaction have to be (quickly!) reset (**local undo**)

# Three Error Categories

## **local error** (in a transaction)

- only one transaction is affected
- modifications caused by this transaction have to be (quickly!) reset (**local undo**)

## **error with loss of DBMS-buffer**

- modifications caused by finished transactions have to be retained (**global redo**)
- modifications caused by unfinished transactions have to be reset (**global undo**).

# Three Error Categories

## **local error** (in a transaction)

- only one transaction is affected
- modifications caused by this transaction have to be (quickly!) reset (**local undo**)

## **error with loss of DBMS-buffer**

- modifications caused by finished transactions have to be retained (**global redo**)
- modifications caused by unfinished transactions have to be reset (**global undo**).

## **error with background memory loss**

- data recovery via backup

# Three Error Categories

## **local error** (in a transaction)

- only one transaction is affected
- modifications caused by this transaction have to be (quickly!) reset (**local undo**)

## **error with loss of DBMS-buffer**

- modifications caused by finished transactions have to be retained (**global redo**)
- modifications caused by unfinished transactions have to be reset (**global undo**).

## **error with background memory loss**

- data recovery via backup

# Impact of Strategy for Outsourcing to Recovery

# Impact of Strategy for Outsourcing to Recovery

	force	$\neg$ force
$\neg$ steal		
steal		



# Impact of Strategy for Outsourcing to Recovery

	force	$\neg$ force
$\neg$ steal	<ul style="list-style-type: none"><li>■ no redo</li><li>■ no undo</li></ul>	
steal		

# Impact of Strategy for Outsourcing to Recovery

	force	$\neg$ force
$\neg$ steal	<ul style="list-style-type: none"><li>■ no redo</li><li>■ no undo</li></ul>	<ul style="list-style-type: none"><li>■ redo</li><li>■ no undo</li></ul>
steal		

# Impact of Strategy for Outsourcing to Recovery

	force	$\neg$ force
$\neg$ steal	<ul style="list-style-type: none"><li>■ no redo</li><li>■ no undo</li></ul>	<ul style="list-style-type: none"><li>■ redo</li><li>■ no undo</li></ul>
steal	<ul style="list-style-type: none"><li>■ no redo</li><li>■ undo</li></ul>	

# Impact of Strategy for Outsourcing to Recovery

	force	$\neg$ force
$\neg$ steal	<ul style="list-style-type: none"><li>■ no redo</li><li>■ no undo</li></ul>	<ul style="list-style-type: none"><li>■ redo</li><li>■ no undo</li></ul>
steal	<ul style="list-style-type: none"><li>■ no redo</li><li>■ undo</li></ul>	<ul style="list-style-type: none"><li>■ redo</li><li>■ undo</li></ul>

# Some Disadvantages of $\neg$ steal + force

- forced outsourcing at transaction end is expensive
  - pages used by many transactions ( “hot spots” ) are outsourced although they remain in the buffer
  - outsourcing of pages has to be atomic ( “everything or nothing” )
- for  $\neg$ steal + force transactions have to block whole page

# Overview

1. Memory Management
2. Error Categories
3. Taken System Configuration
4. Logging
5. Recovery after an Error
6. Checkpoints

# Taken System Configuration

*steal* pages modified by an unfinished transaction can be outsourced as well

→ *force* after transaction end modified pages do not have to be outsourced

*update-in-place* every page has exactly one corresponding place (a copy) in the background memory

*small barrier granulates* on dataset level (see next chapter “concurrency control”): different transactions access page simultaneously; page may contain modifications from both, finished and unfinished transactions

# Overview

1. Memory Management
2. Error Categories
3. Taken System Configuration
- 4. Logging**
  - 4.1 Structure of Log-Entries
  - 4.2 Writing of Log-Entries
5. Recovery after an Error
6. Checkpoints



# Overview

1. Memory Management
2. Error Categories
3. Taken System Configuration
- 4. Logging**
  - 4.1 Structure of Log-Entries**
  - 4.2 Writing of Log-Entries
5. Recovery after an Error
6. Checkpoints

# Protocol of Modifications

## problem:

- **global undo**: background memory may receive modifications of unfinished transactions (because of steal)
- **global redo**: modifications due to finished transactions may not be stored in background memory yet (because of  $\neg$ force)

# Protocol of Modifications

## problem:

- **global undo**: background memory may receive modifications of unfinished transactions (because of steal)
- **global redo**: modifications due to finished transactions may not be stored in background memory yet (because of  $\neg$ force)

## solution: storing additional information (log-file)

- information required from undo/redo
- information about start and end of a transaction

# Structure of Log-Entries

## Definition (structure of a log-entry)

[LSN, transactionID, pageID, redo, undo, prevLSN]

# Structure of Log-Entries

## Definition (structure of a log-entry)

[LSN, transactionID, pageID, redo, undo, prevLSN]

### LSN (Log Sequence Number):

- unique identifier of log-entry
- are assigned monotonically increasing
- chronological order can be investigated

# Structure of Log-Entries

## Definition (structure of a log-entry)

[LSN, transactionID, pageID, redo, undo, prevLSN]

### LSN (Log Sequence Number):

- unique identifier of log-entry
- are assigned monotonically increasing
- chronological order can be investigated

### transactionID:

- identification of transaction that has lead to modification

# Structure of Log-Entries

## Definition (structure of a log-entry)

[LSN, transactionID, pageID, redo, undo, prevLSN]

### LSN (Log Sequence Number):

- unique identifier of log-entry
- are assigned monotonically increasing
- chronological order can be investigated

### transactionID:

- identification of transaction that has lead to modification

### pageID:

- identification of the corresponding page
- one log-entry is created per page affected

# Structure of Log-Entries

## Definition (structure of a log-entry)

[LSN, transactionID, pageID, redo, undo, prevLSN]

redo:

- information to track modification



# Structure of Log-Entries

## Definition (structure of a log-entry)

[LSN, transactionID, pageID, redo, undo, prevLSN]

redo:

- information to track modification

undo:

- information to undo modification

# Structure of Log-Entries

## Definition (structure of a log-entry)

[LSN, transactionID, pageID, redo, undo, prevLSN]

redo:

- information to track modification

undo:

- information to undo modification

prevLSN:

- pointer to previous log-entry of the transaction
- necessary for efficiency reasons (local undo!)

# Example of Log-Entries

## Example

step	$T_1$	$T_2$	Log
1.	<b>BOT</b>		$[\#1, T_1, \text{BOT}, 0]$
2.	$r(A, a_1)$		
3.		<b>BOT</b>	$[\#2, T_2, \text{BOT}, 0]$
4.		$r(C, c_2)$	
6.	$w(A, a_1 - 50)$		$[\#3, T_1, P_A, A-50, A+=50, \#1]$
8.		$w(C, c_2 + 100)$	$[\#4, T_2, P_C, C+=100, C-=100, \#2]$
9.	$r(B, b_1)$		
11.	$w(B, b_1 + 50)$		$[\#5, T_1, P_B, B+=50, B-=50, \#3]$
12.	<b>commit</b>		$[\#6, T_1, \text{commit}, \#5]$
13.		$r(A, a_2)$	
15.		$w(A, a_2 - 100)$	$[\#7, T_2, P_A, A-=100, A+=100, \#4]$
16.		<b>commit</b>	$[\#8, T_2, \text{commit}, \#7]$

# Physical or Logical Logging

physical logging: redo- and undo-entries get state:

- undo receives **before-image** (= state before modification)
- redo receives **after-image** (= state after modification)

# Physical or Logical Logging

**physical logging:** redo- and undo-entries get state:

- undo receives **before-image** (= state before modification)
- redo receives **after-image** (= state after modification)

**logical logging:** redo- and undo-entries get operations:

- undo contains information how **before-image** is constructed from after-image
- redo contains information how **after-image** is constructed from before-image

# Physical versus. Logical Logging

## Example

	$T_1$	$T_2$
1	BOT	
2		BOT
3	read( $A, a_1$ )	
4	write( $A, a_1 + 5$ )	
5		read( $A, a_2$ )
6		write( $A, a_2 + 10$ )
7	abort	

resetting only possible with logical logging

remark: violates isolation (see Concurrency Control)

# Role of LSN (Log Sequence Number)

**problem:** at recovery the DBMS has to check whether the consequences of a log-entry are already contained in a page

# Role of LSN (Log Sequence Number)

**problem:** at recovery the DBMS has to check whether the consequences of a log-entry are already contained in a page

**solution:** each page has a field for a LSN:


- **construction** of a log-entry: copy entry LSN to the page
- page contains **LSN of the most recent log-entry** corresponding to the page



# Overview

1. Memory Management
2. Error Categories
3. Taken System Configuration
- 4. Logging**
  - 4.1 Structure of Log-Entries
  - 4.2 Writing of Log-Entries**
5. Recovery after an Error
6. Checkpoints

# Writing of Log-Entries



# Writing of Log-Entries

- log-buffer for log-entries
  - separated from buffer for pages
- common: ring buffer
  - continuous writing avoids peeks
- log-information is written twice:
  - 1 in a **temporal log** (background memory) for quick access
  - 2 in a **log-repository** for recovery after error with background memory loss

# The WAL-Principle (Write Ahead Log)

the taken configuration *steal*,  $\neg$ *force*, *update-in-place* requires the **WAL-principle** (*write ahead log*):

- before writing a transaction (**successful commit**) all log-entries of the transaction have to be written out (for **redo**)
- before outsourcing a **modified page** all corresponding log-entries have to be written out (for **undo**)
- the **chronological order** has to be preserved during writing out

# Overview

1. Memory Management

2. Error Categories

3. Taken System Configuration

4. Logging

5. Recovery after an Error

5.1 Recovery after an Error with Loss of Buffer

5.2 Recovery after Local Error

5.3 Recovery after Error with Background Memory Loss

6. Checkpoints

# Overview

1. Memory Management

2. Error Categories

3. Taken System Configuration

4. Logging

5. Recovery after an Error


5.1 Recovery after an Error with Loss of Buffer

5.2 Recovery after Local Error


5.3 Recovery after Error with Background Memory Loss

6. Checkpoints

# We Distinguish Two Kinds of Transactions




# We Distinguish Two Kinds of Transactions



- transactions of type  $T_1$  (**winner**):  
effect has to be completely **reconstructed**




# We Distinguish Two Kinds of Transactions




- transactions of type  $T_1$  (**winner**):  
effect has to be completely **reconstructed**
- transactions of type  $T_2$  (**loser**):  
effect has to be completely **revoked**


# Three Stages of Recovery



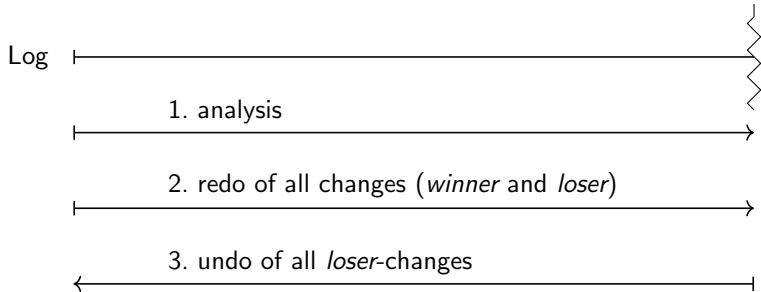
# Three Stages of Recovery




# Three Stages of Recovery



# Three Stages of Recovery




# Three Stages of Recovery: 1. Analysis




- analysis of the whole log-file

# Three Stages of Recovery: 1. Analysis



- analysis of the whole log-file
- identify loser-set  
( = transactions without completion)
- find the highest LSN for each of those transactions


## Three Stages of Recovery: 2. Redo



database is set to state of crash




## Three Stages of Recovery: 2. Redo



database is set to state of crash


- **all** logged modifications are reconstructed
- in **chronological order** (i.e. ordered by LSN)
- **page-LSN** determines state of background memory:
  - LSN of log-entry  $>$  page-LSN: execute redo-operation; update of page-LSN
  - otherwise: background memory contains *after-image* of operation

## Three Stages of Recovery: 3. Undo



removes effects of unfinished transactions


## Three Stages of Recovery: 3. Undo



removes effects of unfinished transactions

- for every log-entry of a loser-transaction: use undo-information for recovery of *before-image*
- in *reverse chronological* order (by LSN)

## Three Stages of Recovery: 3. Undo



removes effects of unfinished transactions

- for every log-entry of a loser-transaction: use undo-information for recovery of *before-image*
- in *reverse chronological* order (by LSN)
- construct *special log-entry* for every executed undo

# Fault Tolerance (Idempotence) of Recovery

**problem:** during recovery an error with loss of buffer may occur

# Fault Tolerance (Idempotence) of Recovery

**problem:** during recovery an error with loss of buffer may occur

**requirements:**

$$\text{undo}(\text{undo}(\dots(\text{undo}(a))\dots)) = \text{undo}(a)$$

$$\text{redo}(\text{redo}(\dots(\text{redo}(a))\dots)) = \text{redo}(a)$$

# Fault Tolerance (Idempotence) of Recovery

**problem:** during recovery an error with loss of buffer may occur

**requirements:**

$$\text{undo}(\text{undo}(\dots(\text{undo}(a))\dots)) = \text{undo}(a)$$

$$\text{redo}(\text{redo}(\dots(\text{redo}(a))\dots)) = \text{redo}(a)$$

**realization:**

- redo-stage: via LSN (in log-entry resp. page-LSN)

# Fault Tolerance (Idempotence) of Recovery

**problem:** during recovery an error with loss of buffer may occur

**requirements:**

$$\text{undo}(\text{undo}(\dots(\text{undo}(a))\dots)) = \text{undo}(a)$$

$$\text{redo}(\text{redo}(\dots(\text{redo}(a))\dots)) = \text{redo}(a)$$

**realization:**

- redo-stage: via LSN (in log-entry resp. page-LSN)
- undo-stage: via **CLRs (compensation log record)**



# Compensation Log Record (CLR)

## Definition (structure of compensation log entry)

$\langle \text{LSN}, \text{transaktionsID}, \text{pageID}, \text{redo}, \text{prevLSN}, \text{UndoNxtLSN} \rangle$

no undo-information:

- revoking of undo-steps **not necessary**:  
either successfully outsourced or *redo*

prevLSN (similar to log-entries):

- Pointer to previous log- or *CLR-entry* of the transaction


UndoNxtLSN:

- pointer to modification of transaction that has to be **revoked next**
- **prevLSN** of revoked log-entry


## Example: Compensation entries in Log



## Example: Compensation entries in Log

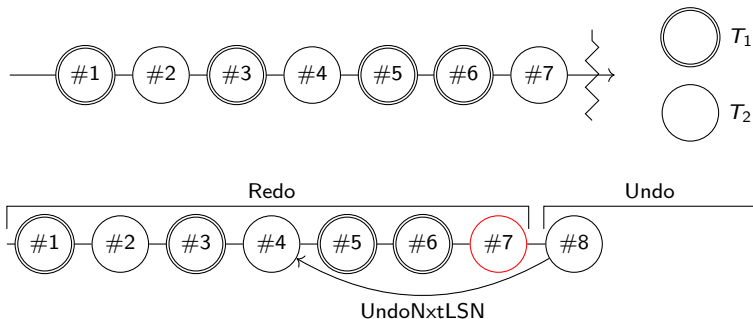


## Example: Compensation entries in Log



CLRs for revoked changes and for BOT:

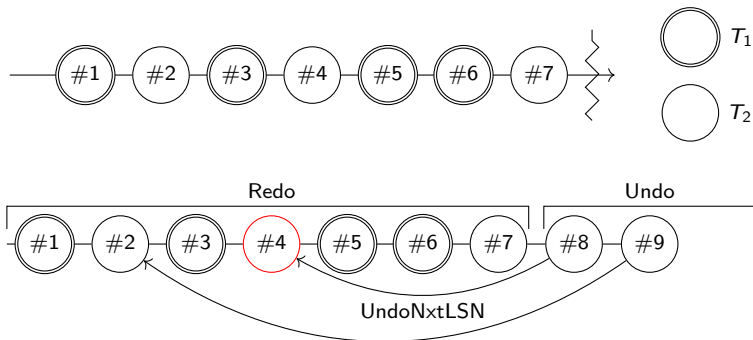
## Example: Compensation entries in Log



CLRs for revoked changes and for BOT:

- #8 is CLR for #7

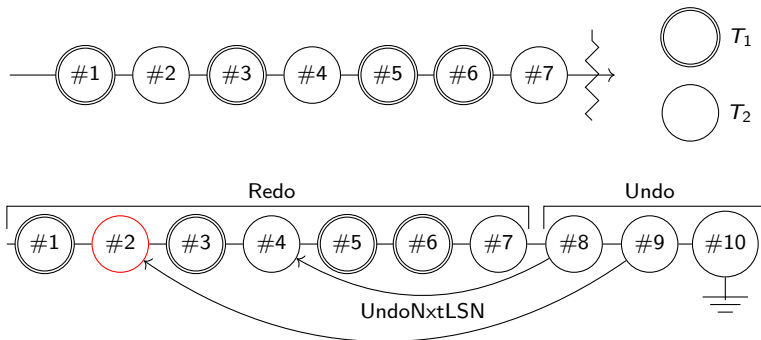
## Example: Compensationentries in Log



CLRs for revoked changes and for BOT:

- #8 is CLR for #7
- #9 is CLR for #4

## Example: Compensationentries in Log



CLRs for revoked changes and for BOT:

- #8 is CLR for #7
- #9 is CLR for #4
- #10 is CLR for #2 (= BOT)

# Example Log-Entries and Recovery

## Example



# Example Log-Entries and Recovery

## Example

$[\#1, T_1, \text{BOT}, 0]$

# Example Log-Entries and Recovery

## Example

$[\#1, T_1, \text{BOT}, 0]$

$[\#2, T_2, \text{BOT}, 0]$

# Example Log-Entries and Recovery

## Example

$[ \#1, T_1, \text{BOT}, 0 ]$   
 $[ \#2, T_2, \text{BOT}, 0 ]$   
 $[ \#3, T_1, P_A, A- = 50, A+ = 50, \#1 ]$

# Example Log-Entries and Recovery

## Example

$[ \#1, T_1, \text{BOT}, 0 ]$   
 $[ \#2, T_2, \text{BOT}, 0 ]$   
 $[ \#3, T_1, P_A, A- = 50, A+ = 50, \#1 ]$   
 $[ \#4, T_2, P_C, C+ = 100, C- = 100, \#2 ]$

# Example Log-Entries and Recovery

## Example

$[ \#1, T_1, \text{BOT}, 0 ]$   
 $[ \#2, T_2, \text{BOT}, 0 ]$   
 $[ \#3, T_1, P_A, A- = 50, A+ = 50, \#1 ]$   
 $[ \#4, T_2, P_C, C+ = 100, C- = 100, \#2 ]$   
 $[ \#5, T_1, P_B, B+ = 50, B- = 50, \#3 ]$

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A=100$ ,  $A+=100$ , #4]

# Example Log-Entries and Recovery

## Example

$[#1, T_1, \text{BOT}, 0]$   
 $[#2, T_2, \text{BOT}, 0]$   
 $[#3, T_1, P_A, A-=50, A+=50, \#1]$   
 $[#4, T_2, P_C, C+=100, C-=100, \#2]$   
 $[#5, T_1, P_B, B+=50, B-=50, \#3]$   
 $[#6, T_1, \text{commit}, \#5]$   
 $[#7, T_2, P_A, A-=100, A+=100, \#4]$



# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A=100$ ,  $A+=100$ , #4]

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A=100$ ,  $A+=100$ , #4]

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A=100$ ,  $A+=100$ , #4]

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A=100$ ,  $A+=100$ , #4]

## Example Log-Entries and Recovery

### Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]  
⟨#8,  $T_2$ ,  $P_A$ ,  $A+=100$ , #7, #4⟩



# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]  
⟨#8,  $T_2$ ,  $P_A$ ,  $A+=100$ , #7, #4⟩  
⟨#9,  $T_2$ ,  $P_C$ ,  $C-=100$ , #8, #2⟩

# Example Log-Entries and Recovery

## Example

[#1,  $T_1$ , BOT, 0]  
[#2,  $T_2$ , BOT, 0]  
[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]  
[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]  
[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]  
[#6,  $T_1$ , commit, #5]  
[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]  
⟨#8,  $T_2$ ,  $P_A$ ,  $A+=100$ , #7, #4⟩  
⟨#9,  $T_2$ ,  $P_C$ ,  $C-=100$ , #8, #2⟩  
⟨#10,  $T_2$ , -, -, #9, 0⟩

# Recovery with CLR

## 2. redo:

- edit whole log and CLRs

## 3. undo:

- manage list of all losers + highest current LSN per loser
- for the loser with maximal current LSN:
  - “normal” log-entry:
    - (1) undo; (2) construct CLR;
    - (3) set current LSN to PrevLSN
  - CLR:
    - (1) set current LSN to UndoNxtLSN

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$   
 $r_1(A, a_1)$   
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$   
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 $\dots$

## Log Entries

## Pages

$P_A$	LSN: #0
$A = 75$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #0
$C = 10$	

**active trans.**

$Tr.$	LSN

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$   
 $r_1(A, a_1)$   
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$   
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 $\dots$

## Log Entries

$[ \#1, T_3, BOT, \#0 ]$   
 $[ \#2, T_1, BOT, \#0 ]$   
 $[ \#3, T_2, BOT, \#0 ]$

## Pages

$P_A$	LSN: #0
$A = 75$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #0
$C = 10$	

## active trans.

$Tr.$	LSN
1	2
2	3
3	1

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$  [10]  
 $r_1(A, a_1)$   
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$   
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 $\dots$

## Log Entries

$[ \#1, T_3, BOT, \#0 ]$   
 $[ \#2, T_1, BOT, \#0 ]$   
 $[ \#3, T_2, BOT, \#0 ]$

## Pages

$P_A$	LSN: #0
$A = 75$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #0
$C = 10$	

## active trans.

$Tr.$	LSN
1	2
2	3
3	1

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$  [10]  
 $r_1(A, a_1)$  [75]  
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$   
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 $\dots$

## Log Entries

$[ \#1, T_3, BOT, \#0 ]$   
 $[ \#2, T_1, BOT, \#0 ]$   
 $[ \#3, T_2, BOT, \#0 ]$

## Pages

$P_A$	LSN: #0
$A = 75$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #0
$C = 10$	

## active trans.

$Tr.$	LSN
1	2
2	3
3	1

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$  [10]  
 $r_1(A, a_1)$  [75]  
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$   
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 $\dots$

## Log Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

## Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #0
$C = 10$	

## active trans.

$Tr.$	LSN
1	4
2	3
3	1



# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$  [10]  
 $r_1(A, a_1)$  [75]  
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$  [120]  
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 ...

## Log Entries

$[#1, T_3, BOT, \#0]$   
 $[#2, T_1, BOT, \#0]$   
 $[#3, T_2, BOT, \#0]$   
 $[#4, T_1, P_A, A-=50, A+=50, \#2]$

## Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #0
$C = 10$	

## active trans.

$Tr.$	LSN
1	4
2	3
3	1

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$  [10]  
 $r_1(A, a_1)$  [75]  
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$  [120]  
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 ...

## Log Entries

$[#1, T_3, BOT, \#0]$   
 $[#2, T_1, BOT, \#0]$   
 $[#3, T_2, BOT, \#0]$   
 $[#4, T_1, P_A, A-=50, A+=50, \#2]$   
 $[#5, T_2, P_A, A+=0, A-=0, \#3]$

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #0
$C = 10$	

## active trans.

$Tr.$	LSN
1	4
2	5
3	1

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$  [10]  
 $r_1(A, a_1)$  [75]  
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$  [120]  
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 ...

## Log Entries

$[ \#1, T_3, BOT, \#0 ]$   
 $[ \#2, T_1, BOT, \#0 ]$   
 $[ \#3, T_2, BOT, \#0 ]$   
 $[ \#4, T_1, P_A, A-=50, A+=50, \#2 ]$   
 $[ \#5, T_2, P_A, A+=0, A-=0, \#3 ]$   
 $[ \#6, T_3, P_C, C+=25, C-=25, \#1 ]$

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #6
$C = 35$	

## active trans.

$Tr.$	LSN
1	4
2	5
3	6

# Example: Logging and Recovery

## Schedule

$BOT_3$   
 $BOT_1$   
 $BOT_2$   
 $r_3(C, c_3)$  [10]  
 $r_1(A, a_1)$  [75]  
 $w_1(A, a_1 - 50)$   
 $r_2(B, b_2)$  [120]  
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$   
 $w_1(C, c_1 + 25)$   
 ...

## Log Entries

$[#1, T_3, BOT, \#0]$   
 $[#2, T_1, BOT, \#0]$   
 $[#3, T_2, BOT, \#0]$   
 $[#4, T_1, P_A, A-=50, A+=50, \#2]$   
 $[#5, T_2, P_A, A+=0, A-=0, \#3]$   
 $[#6, T_3, P_C, C+=25, C-=25, \#1]$   
 $[#7, T_2, P_B, B+=50, B-=50, \#5]$

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #7
$B = 170$	

$P_C$	LSN: #6
$C = 35$	

## active trans.

Tr.	LSN
1	4
2	7
3	6

# Example: Logging and Recovery

## Schedule

$BOT_3$   $[#1, T_3, BOT, \#0]$   
 $BOT_1$   $[#2, T_1, BOT, \#0]$   
 $BOT_2$   $[#3, T_2, BOT, \#0]$   
 $r_3(C, c_3)$  [10]  $[#4, T_1, P_A, A-=50, A+=50, \#2]$   
 $r_1(A, a_1)$  [75]  $[#5, T_2, P_A, A+=0, A-=0, \#3]$   
 $w_1(A, a_1 - 50)$   $[#6, T_3, P_C, C+=25, C-=25, \#1]$   
 $r_2(B, b_2)$  [120]  $[#7, T_2, P_B, B+=50, B-=50, \#5]$   
 $w_2(A, b_2 - 95)$   
 $w_3(C, c_3 + 25)$   
 $w_2(B, b_2 + 50)$   
 $r_1(C, c_1)$  [35]  
 $w_1(C, c_1 + 25)$   
 ...

## Log Entries

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #7
$B = 170$	

$P_C$	LSN: #6
$C = 35$	

## active trans.

$Tr.$	LSN
1	4
2	7
3	6

# Example: Logging and Recovery

## Schedule

$BOT_3$

$BOT_1$

$BOT_2$

$r_3(C, c_3)$  [10]

$r_1(A, a_1)$  [75]

$w_1(A, a_1 - 50)$

$r_2(B, b_2)$  [120]

$w_2(A, b_2 - 95)$

$w_3(C, c_3 + 25)$

$w_2(B, b_2 + 50)$

$r_1(C, c_1)$  [35]

$w_1(C, c_1 + 25)$

...

## Log Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A- = 50$ ,  $A+ = 50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+ = 0$ ,  $A- = 0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+ = 25$ ,  $C- = 25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+ = 50$ ,  $B- = 50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+ = 25$ ,  $C- = 25$ , #4]

## Pages

$P_A$  LSN: #5

$A = 25$

$P_B$  LSN: #7

$B = 170$

$P_C$  LSN: #8

$C = 60$

## active trans.

Tr.	LSN
1	8
2	7
3	6

# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 $\text{commit}_3$   
 $\text{abort}_2$

## Log Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #7
$B = 170$	

$P_C$	LSN: #8
$C = 60$	

## active trans.

$Tr.$	LSN
1	8
2	7
3	6

# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 $\text{commit}_3$   
 $\text{abort}_2$

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

## Log Entries

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #8
$C = 60$	

## active trans.

Tr.	LSN
1	8
2	7
3	9



# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 commit<sub>3</sub>  
 abort<sub>2</sub>

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

## Log Entries

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #10
$C = 85$	

## active trans.

Tr.	LSN
1	8
2	10
3	9

# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 commit<sub>3</sub>  
 abort<sub>2</sub>

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]

## Log Entries

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #10
$C = 85$	

## active trans.

Tr.	LSN
1	8
2	10

# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 commit<sub>3</sub>  
 abort<sub>2</sub>

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #10
$C = 85$	

## active trans.

Tr.	LSN
1	8
2	10

# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 $\text{commit}_3$   
 $\text{abort}_2$

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #10
$C = 85$	

## active trans.

Tr.	LSN
1	8
2	10

# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 commit<sub>3</sub>  
 abort<sub>2</sub>  
 [#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 <[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7>

## Log Entries

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

## active trans.

Tr.	LSN
1	8
2	7

# Example: Logging and Recovery

## Schedule

...  
 $r_2(C, c_2)$  [60]  
 $r_3(B, b_3)$  [170]  
 $w_3(B, b_3 - 75)$   
 $w_2(C, c_2 + 25)$   
 commit<sub>3</sub>  
 abort<sub>2</sub>  
 [#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ , A-=50, A+=50, #2]  
 [#5,  $T_2$ ,  $P_A$ , A+=0, A-=0, #3]  
 [#6,  $T_3$ ,  $P_C$ , C+=25, C-=25, #1]  
 [#7,  $T_2$ ,  $P_B$ , B+=50, B-=50, #5]  
 [#8,  $T_1$ ,  $P_C$ , C+=25, C-=25, #4]  
 [#9,  $T_3$ ,  $P_B$ , B-=75, B+=75, #6]  
 [#10,  $T_2$ ,  $P_C$ , C+=25, C-=25, #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 <#12,  $T_2$ ,  $P_C$ , C-=25, #10, #7>

crash

## Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

## active trans.

Tr.	LSN
1	8
2	7

# Example: Logging and Recovery

## Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

## Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

# Example: Logging and Recovery

## 1. analysis stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	Status	
1	loser	
2	loser	
3	winner	



# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	2	
2	3	
3	1	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	2	
2	3	
3	1	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	2	
2	3	
3	1	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	4	
2	3	
3	1	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	4	
2	3	
3	1	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #4
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN
1	4
2	3
3	1

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN
1	4
2	5
3	1

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

- [#1,  $T_3$ , BOT, #0]
- [#2,  $T_1$ , BOT, #0]
- [#3,  $T_2$ , BOT, #0]
- [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓
- [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗
- [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓
- [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]
- [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]
- [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]
- [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]
- [#11,  $T_3$ , COMMIT, #9]
- ⟨#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7⟩

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #0
$B = 120$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	4	
2	5	
3	6	



# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

- [#1,  $T_3$ , BOT, #0]
- [#2,  $T_1$ , BOT, #0]
- [#3,  $T_2$ , BOT, #0]
- [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓
- [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗
- [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓
- [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5] ✗
- [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]
- [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]
- [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]
- [#11,  $T_3$ , COMMIT, #9]
- ⟨#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7⟩

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #7
$B = 170$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	4	
2	7	
3	6	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

- [#1,  $T_3$ , BOT, #0]
- [#2,  $T_1$ , BOT, #0]
- [#3,  $T_2$ , BOT, #0]
- [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓
- [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗
- [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓
- [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5] ✗
- [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4] ✓
- [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]
- [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]
- [#11,  $T_3$ , COMMIT, #9]
- ⟨#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7⟩

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #7
$B = 170$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	8	
2	7	
3	6	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

- [#1,  $T_3$ , BOT, #0]
- [#2,  $T_1$ , BOT, #0]
- [#3,  $T_2$ , BOT, #0]
- [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓
- [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗
- [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓
- [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5] ✗
- [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4] ✓
- [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6] ✗
- [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]
- [#11,  $T_3$ , COMMIT, #9]
- ⟨#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7⟩

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #8
$C = 60$	

### transactions

Tr.	LSN	
1	8	
2	7	
3	9	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5] ✗  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4] ✓  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6] ✗  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7] ✗  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #10
$C = 85$	

### transactions

Tr.	LSN	
1	8	
2	10	
3	9	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5] ✗  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4] ✓  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6] ✗  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7] ✗  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #10
$C = 85$	

### transactions

Tr.	LSN	
1	8	
2	10	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

- [#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5] ✗  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4] ✓  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6] ✗  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7] ✗  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7] ✗

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

### transactions

Tr.	LSN	
1	8	
2	12	

# Example: Logging and Recovery

## 2. redo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2] ✓  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3] ✗  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1] ✓  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5] ✗  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4] ✓  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6] ✗  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7] ✗  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7] ✗

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

### transactions

Tr.	LSN	
1	8	
2	12	

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

### transactions

$Tr.$	NxtU	LSN
1	8	8
2	12	12



# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

### transactions

Tr.	NxtU	LSN
1	8	8
2	12	12

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

### transactions

$Tr.$	NxtU	LSN
1	8	8
2	7	12

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

### transactions

$Tr.$	NxtU	LSN
1	8	8
2	7	12

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #12
$C = 60$	

### transactions

$Tr.$	NxtU	LSN
1	8	8
2	7	12

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

⟨#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7⟩

→ ⟨#13,  $T_1$ ,  $P_C$ ,  $C-=25$ , #8, #4⟩

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #13
$C = 35$	

### transactions

Tr.	NxtU	LSN
1	4	13
2	7	12

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #9
$B = 95$	

$P_C$	LSN: #13
$C = 35$	

### transactions

$Tr.$	NxtU	LSN
1	4	13
2	7	12

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]  
 → [#14,  $T_2$ ,  $P_B$ ,  $B-=50$ , #12, #5]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

Tr.	NxtU	LSN
1	4	13
2	5	14

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #5
$A = 25$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

Tr.	NxtU	LSN
1	4	13
2	5	14



# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]  
 → [#15,  $T_2$ ,  $P_A$ ,  $A-=0$ , #14, #3]

### Pages

$P_A$	LSN: #15
$A = 25$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

Tr.	NxtU	LSN
1	4	13
2	3	15

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]  
 [#2,  $T_1$ , BOT, #0]  
 [#3,  $T_2$ , BOT, #0]  
 [#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]  
 [#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]  
 [#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]  
 [#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]  
 [#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]  
 [#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]  
 [#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]  
 [#11,  $T_3$ , COMMIT, #9]  
 [#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #15
$A = 25$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

$Tr.$	NxtU	LSN
1	4	13
2	3	15

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

→ [#16,  $T_1$ ,  $P_A$ ,  $A+=50$ , #13, #2]

### Pages

$P_A$	LSN: #16
$A = 75$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

Tr.	NxtU	LSN
1	2	16
2	3	15

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #16
$A = 75$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

$Tr.$	NxtU	LSN
1	2	16
2	3	15

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

⟨#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7⟩

→ ⟨#17,  $T_2$ , BOT, #15⟩

### Pages

$P_A$	LSN: #16
$A = 75$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

$Tr.$	NxtU	LSN
1	2	16

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

[#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7]

### Pages

$P_A$	LSN: #16
$A = 75$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

Tr.	NxtU	LSN
1	2	16

# Example: Logging and Recovery

## 3. undo-stage

### Log-Entries

[#1,  $T_3$ , BOT, #0]

[#2,  $T_1$ , BOT, #0]

[#3,  $T_2$ , BOT, #0]

[#4,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #2]

[#5,  $T_2$ ,  $P_A$ ,  $A+=0$ ,  $A-=0$ , #3]

[#6,  $T_3$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #1]

[#7,  $T_2$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #5]

[#8,  $T_1$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #4]

[#9,  $T_3$ ,  $P_B$ ,  $B-=75$ ,  $B+=75$ , #6]

[#10,  $T_2$ ,  $P_C$ ,  $C+=25$ ,  $C-=25$ , #7]

[#11,  $T_3$ , COMMIT, #9]

⟨#12,  $T_2$ ,  $P_C$ ,  $C-=25$ , #10, #7⟩

→ ⟨#18,  $T_1$ , BOT, #16⟩

### Pages

$P_A$	LSN: #16
$A = 75$	

$P_B$	LSN: #14
$B = 45$	

$P_C$	LSN: #13
$C = 35$	

### transactions

Tr.	NxtU	LSN

# Recovery with CLR – Example

## Example

$[\#1, T_1, \text{BOT}, 0]$

$[\#2, T_2, \text{BOT}, 0]$

$[\#3, T_1, P_A, A-50, A+=50, \#1]$

$[\#4, T_2, P_C, C+=100, C-=100, \#2]$

$[\#5, T_1, P_B, B+=50, B-=50, \#3]$

$[\#6, T_1, \text{commit}, \#5]$

$[\#7, T_2, P_A, A-=100, A+=100, \#4]$

$\langle \#8, T_2, P_A, A+=100, \#7, \#4 \rangle$



# Recovery with CLR – Example

## Example

[#1,  $T_1$ , BOT, 0]

[#2,  $T_2$ , BOT, 0]

[#3,  $T_1$ ,  $P_A$ ,  $A- = 50$ ,  $A+ = 50$ , #1]

[#4,  $T_2$ ,  $P_C$ ,  $C+ = 100$ ,  $C- = 100$ , #2]

[#5,  $T_1$ ,  $P_B$ ,  $B+ = 50$ ,  $B- = 50$ , #3]

[#6,  $T_1$ , commit, #5]

[#7,  $T_2$ ,  $P_A$ ,  $A- = 100$ ,  $A+ = 100$ , #4]

$\langle$  #8,  $T_2$ ,  $P_A$ ,  $A+ = 100$ , #7, #4  $\rangle$

### redo-phase:

- log-entries #1 – #8
- consider page-LSN

# Recovery with CLR – Example

## Example

[#1,  $T_1$ , BOT, 0]

[#2,  $T_2$ , BOT, 0]

[#3,  $T_1$ ,  $P_A$ ,  $A-50$ ,  $A+=50$ , #1]

[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]

[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]

[#6,  $T_1$ , commit, #5]

[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

$\langle \#8, T_2, P_A, A+=100, \#7, \#4 \rangle$

### redo-phase:

- log-entries #1 – #8
- consider page-LSN

### undo-phase:

- follow UndoNxtLSN

# Recovery with CLR – Example

## Example

[#1,  $T_1$ , BOT, 0]

[#2,  $T_2$ , BOT, 0]

[#3,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #1]

[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]

[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]

[#6,  $T_1$ , commit, #5]

[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

⟨#8,  $T_2$ ,  $P_A$ ,  $A+=100$ , #7, #4⟩

### redo-phase:

- log-entries #1 – #8
- consider page-LSN

### undo-phase:

- follow UndoNxtLSN

# Recovery with CLR – Example

## Example

[#1,  $T_1$ , BOT, 0]

[#2,  $T_2$ , BOT, 0]

[#3,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #1]

[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]

[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]

[#6,  $T_1$ , commit, #5]

[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

⟨#8,  $T_2$ ,  $P_A$ ,  $A+=100$ , #7, #4⟩

### redo-phase:

- log-entries #1 – #8
- consider page-LSN

### undo-phase:

- follow UndoNxtLSN

# Recovery with CLR – Example

## Example

$[ \#1, T_1, \text{BOT}, 0 ]$

$[ \#2, T_2, \text{BOT}, 0 ]$

$[ \#3, T_1, P_A, A-50, A+=50, \#1 ]$

$[ \#4, T_2, P_C, C+=100, C-=100, \#2 ]$

$[ \#5, T_1, P_B, B+=50, B-=50, \#3 ]$

$[ \#6, T_1, \text{commit}, \#5 ]$

$[ \#7, T_2, P_A, A-=100, A+=100, \#4 ]$

$\langle \#8, T_2, P_A, A+=100, \#7, \#4 \rangle$

$\langle \#9, T_2, P_C, C-=100, \#8, \#2 \rangle$

$\langle \#10, T_2, -, -, \#9, 0 \rangle$

# Recovery with CLR – Example

## Example

[#1,  $T_1$ , BOT, 0]

[#2,  $T_2$ , BOT, 0]

[#3,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #1]

[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]

[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]

[#6,  $T_1$ , commit, #5]

[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

$\langle \#8, T_2, P_A, A+=100, \#7, \#4 \rangle$

$\langle \#9, T_2, P_C, C-=100, \#8, \#2 \rangle$

$\langle \#10, T_2, -, -, \#9, 0 \rangle$

### redo-phase:

- log entries #1 – #10
- note the page-LSN

# Recovery with CLR – Example

## Example

[#1,  $T_1$ , BOT, 0]

[#2,  $T_2$ , BOT, 0]

[#3,  $T_1$ ,  $P_A$ ,  $A-=50$ ,  $A+=50$ , #1]

[#4,  $T_2$ ,  $P_C$ ,  $C+=100$ ,  $C-=100$ , #2]

[#5,  $T_1$ ,  $P_B$ ,  $B+=50$ ,  $B-=50$ , #3]

[#6,  $T_1$ , commit, #5]

[#7,  $T_2$ ,  $P_A$ ,  $A-=100$ ,  $A+=100$ , #4]

$\langle$  #8,  $T_2$ ,  $P_A$ ,  $A+=100$ , #7, #4  $\rangle$

$\langle$  #9,  $T_2$ ,  $P_C$ ,  $C-=100$ , #8, #2  $\rangle$

$\langle$  #10,  $T_2$ , -, -, #9, 0  $\rangle$

### redo-phase:

- log entries #1 – #10
- note the page-LSN

### undo-phase:

- follow UndoNxtLSN
- UndoNxtLSN = 0  
 $\Rightarrow T_2$  is done

# ARIES

- described procedure forms the kernel of the **ARIES** procedure
- further refinement mainly for **reduction of access** to the **background memory**
- very flexible recovery method



# ARIES

- described procedure forms the kernel of the **ARIES** procedure
- further refinement mainly for **reduction of access** to the **background memory**
- very flexible recovery method
- core components of ARIES:
  - **WAL**
  - recovery of **complete** history in **redo**-step
  - construction of **CLRs** in **undo**-step

# Overview

1. Memory Management

2. Error Categories

3. Taken System Configuration

4. Logging

**5. Recovery after an Error**

5.1 Recovery after an Error with Loss of Buffer

**5.2 Recovery after Local Error**

5.3 Recovery after Error with Background Memory Loss

6. Checkpoints

# Local Reset of Transaction

- 1 investigate the last log-entry of the transaction
- 2 **local undo**: modifications of the transaction are revoked
  - log passed against chronological order
  - CLRs are constructed

# Local Reset of Transaction

- 1 investigate the last log-entry of the transaction
- 2 **local undo**: modifications of the transaction are revoked
  - log passed against chronological order
  - CLRs are constructed

## efficiency:

- usually for every transaction a **pointer to the last log-entry** is maintained
- **backward chaining** via **PrevLSN** allows for quick log traversing
- most **log-entries** of active transactions with reasonable size are still **in the buffer**

# Partial Reset of a Transaction


- 1 reset modifications and construct CLRs starting from the last log-entry (as for a **local undo**).
- 2 **difference**: do not go back to BOT, but reset the transaction up to the intended point

save points possible

# Overview

1. Memory Management
2. Error Categories
3. Taken System Configuration
4. Logging
5. Recovery after an Error
  - 5.1 Recovery after an Error with Loss of Buffer
  - 5.2 Recovery after Local Error
  - 5.3 Recovery after Error with Background Memory Loss
6. Checkpoints

# Media-Recovery



# Overview

1. Memory Management
2. Error Categories
3. Taken System Configuration
4. Logging
5. Recovery after an Error
- 6. Checkpoints**



# Why Checkpoints

**problem** with recovery method:

- we have to traverse the whole log
- log gets bigger and bigger after some time

# Why Checkpoints

**problem** with recovery method:

- we have to traverse the whole log
- log gets bigger and bigger after some time

**solution:**

- enforce (occasional) writing of modified pages = checkpoint
- this way the log is needed only after a certain LSN
- the required minimal LSN is based on time and the kind of the checkpoint

# Why Checkpoints

**problem** with recovery method:

- we have to traverse the whole log
- log gets bigger and bigger after some time

**solution:**

- enforce (occasional) writing of modified pages = checkpoint
- this way the log is needed only after a certain LSN
- the required minimal LSN is based on time and the kind of the checkpoint

**attention!** checkpoint  $\neq$  save point!

# Kinds of Checkpoints

## transaction-consistent checkpoints

- “best” quality for recovery, but expensive


## action-consistent checkpoints

- *undo* requires “older” log-entries, leads to peaks in load

## fuzzy checkpoints

- *undo* and *redo* require “older” log-entries, continuous writing

# Transaction-Consistent Checkpoints



# Transaction-Consistent Checkpoints

## idea:

- background memory should receive **all modifications** of transactions completed at time  $S_i$   
⇒ **no redo beyond  $S_i$**  needed
- at time  $S_i$  **no active transactions** allowed  
⇒ **no undo beyond  $S_i$**  needed

# Transaction-Consistent Checkpoints


## idea:

- background memory should receive **all modifications** of transactions completed at time  $S_i$   
⇒ **no redo beyond  $S_i$**  needed
- at time  $S_i$  **no active transactions** allowed  
⇒ **no undo beyond  $S_i$**  needed

## problem:

- **no new transactions** are allowed to start between signing in and writing of checkpoints
- leads to usually not acceptable latency

# Action-Consistent Checkpoints





# Action-Consistent Checkpoints

## idea:

- all active modifying operations should be terminated
- then writing all modified pages
- only start of next modifying operation is delayed, but not the start of new transactions

# Action-Consistent Checkpoints

## idea:

- all active modifying operations should be terminated
- then writing all modified pages
- only start of next modifying operation is delayed, but not the start of new transactions

## recovery:

- analysis-stage starting from  $S_i$
- no redo beyond  $S_i$  necessary
- in general an undo beyond  $S_i$  is necessary; up to  $MinLSN$  (= smallest LSN of the transaction active at checkpoint)

# Fuzzy Checkpoints


## idea:

- pages should be **written continuously**  
no abrupt writing of many pages because of checkpoint
- at checkpoint only the **identification of all modified pages** (= dirty pages) is written
- additionally *MinDirtyPageLSN* (= minimal LSN whose modifications have not been written yet) is maintained and written at checkpoint


## problem:

- „**hot-spots**“ (constantly needed pages) are not written for a long time, writing is **enforced** when a page occurred among the „dirty pages“ several times


# Three Checkpoint Qualities – Summary




## (a) transaction consistent



## (b) action consistent



## (c) fuzzy



# Learning Objectives

- What are pages?
- Which strategies for the replacement/outsourcing of buffer pages are there? Advantages? Disadvantages?
- Which introducing strategies are there and how do they work?
- Which kinds of error categories are there?
- What kind of information is stored in a log-entry?
  - Difference between logical and physical logging.
  - What is the LSN? Why do we need it?
- What is the WAL-principle?
- How does recovery work after an error with loss of buffer?
  - Stages of recovery and how do they work?
  - How is fault tolerance for recovery implemented?
  - What are CLRs?
- Why do we need checkpoint and which kinds of checkpoint are there?