

# Data Base Systems

VU 184.686, WS 2020

SQL Details

Anela Lolić

Institute of Logic and Computation, TU Wien



FAKULTÄT  
FÜR INFORMATIK

Faculty of Informatics

# Acknowledgements

The slides are based on the slides (in German) of [Sebastian Skritek](#).

The Section on Views is based on [Chapter 4](#) of  
(Kemper, Eickler: Datenbanksysteme – Eine Einführung)

The remaining topics are based on the respective chapters in the  
PostgreSQL Online-Documentation:  
<https://www.postgresql.org/docs/current/static/>

# SQL Details

1. Overview
2. Decomposing Complex Queries and Recursive Queries
3. Views
4. Window Functions
5. Sequences
6. (Other) Data Types and Built-In Functions
7. Learning Objectives

# Topics

1. Overview
2. Decomposing Complex Queries and Recursive Queries
3. Views
4. Window Functions
5. Sequences
6. (Other) Data Types and Built-In Functions
7. Learning Objectives

# Topics

1. Overview
2. Decomposing Complex Queries and Recursive Queries
3. Views
4. Window Functions
5. Sequences
6. (Other) Data Types and Built-In Functions
7. Learning Objectives

# Decomposing Complex Queries

## problem:

- nested SQL queries sometimes become confusing and unclear
- in many cases very similar sub queries

# Decomposing Complex Queries

## problem:

- nested SQL queries sometimes become confusing and unclear
- in many cases very similar sub queries

## Example (nested queries)

find all those students, that have the best average grades

```
SELECT matrNr, avg(grade)
FROM examine
GROUP BY matrNr
HAVING avgGrade >= all (SELECT avg(grade)
                        FROM examine
                        GROUP BY matrNr)
```

# Decomposing Complex Queries

## Example (SQL WITH statement (slight “overkill”))

find all those students, that have the best average grades

```
WITH avgGrade AS (  
    SELECT matrNr, avg(grade) AS dsN  
    FROM examine  
    GROUP BY matrNr ),  
bestavgGrade AS (  
    SELECT matrNr, dsN  
    FROM avgGrade  
    WHERE dsN = (SELECT max(dsN)  
                FROM avgGrade))  
SELECT matrNr, dsN  
FROM bestavgGrade
```



# SQL: SELECT in WITH

WITH queries (also **Common Table Expressions (CTE)**):

- allow for the definition of **temporary “tables”**
- only visible and usable in the current query
- useful for **structuring** the query

# SQL: SELECT in WITH

WITH queries (also **Common Table Expressions (CTE)**):

- allow for the definition of **temporary “tables”**
- only visible and usable in the current query
- useful for **structuring** the query

distinction: WITH and WITH RECURSIVE

**WITH** “syntactic sugar” — no additional expressive power

**WITH RECURSIVE** allows for “recursive” queries — was not possible before

# Recursive Queries

**problem:** SQL queries we have had so far do not allow for querying transitive properties:

## Example

- ancestors or descendants from the parent-child relationship
- all preconditions (and their preconditions, and ...) of a lecture
- all answers or contributions to a discussion
- reachability in graphs
- ...

# Recursive Queries: General Structure

- query that is able to work on its own output

```
WITH RECURSIVE tabName(attrList) AS (  
    non recursive part  
    UNION [ ALL ]  
    redursive part  
    )  
SELECT ...
```

- only the **recursive part** may link to its own output
- both parts contain sub queries (SELECT)
- output has to correspond to *attrList*
- *attrList* contains only names

# Recursive Queries: Semantics (Idea)

- 1 non recursive part constructs a **result set**  $T_0$
- 2 recursive part constructs a **result set**  $T_{i+1}$  from input  $T_i$
- 3 step 2 is repeated until  $T_{i+1} = \emptyset$
- 4 entire result set corresponds to UNION resp. UNION ALL of  $T_0, \dots, T_n$

remark:

- when using **UNION** (instead of UNION ALL):  $T_i \cap \bigcup_{j=0}^{i-1} T_j = \emptyset$
- **attention:** in case the recursive part is generating always (new) tuples: **endless loop**

# Recursive Queries: Semantics

- three tuple sets: result, *working table*, *intermediate table*

# Recursive Queries: Semantics

- three tuple sets: result, *working table*, *intermediate table*

## execution:

- 1 evaluate the non recursive part
  - in case of UNION: eliminate duplicates
  - copy remaining tuples both into the **result set** and in the *working table*

# Recursive Queries: Semantics

- three tuple sets: result, *working table*, *intermediate table*

## execution:

- 2 as long as the *working table* is not empty repeat:
  - 1 evaluate the *recursive part*:
    - *working table* is used as input for the self reference
    - store the result in the *intermediate table*
  - 2 in case of **UNION**: remove all *duplicates* from *intermediate table* and tuples that are already contained in the *result set*
  - 3 *insert content* of *intermediate table* into the result set
  - 4 replace content of the *working table* with the one of the *intermediate table*, then empty *intermediate table*



# Recursive Queries (Example)

Example (all numbers from 1 to 100)

```
WITH RECURSIVE t(n) AS (  
    SELECT 1  
    UNION ALL  
    SELECT n+1 FROM t WHERE n < 100  
)  
SELECT n FROM t;
```

# Recursive Queries (Example)

## Example (all preconditions of lecture 5216)

```
WITH RECURSIVE pred(v) AS (  
    SELECT predNr  
    FROM presuppose  
    WHERE succNr = 5216  
    UNION  
    SELECT presuppose.predNr  
    FROM pred, presuppose  
    WHERE pred.v=presuppose.succNr  
)  
SELECT v FROM pred;
```

# Recursive Queries (Example)

remark: self reference is not permitted in nested queries

Example (all preconditions of lecture 5216)

```
WITH RECURSIVE pred(v) AS (  
    SELECT predNr FROM presuppose  
    WHERE succNr = 5216  
    UNION  
    SELECT predNr  
    FROM presuppose  
    WHERE succNr in (SELECT v FROM pred)  
) SELECT v FROM pred;
```

results in: **ERROR: recursive reference to query**  
**"pred" must not appear within a subquery LINE 8: WHERE**  
**succNr in (SELECT v FROM pred)**

# Topics

## 1. Overview

## 2. Decomposing Complex Queries and Recursive Queries

## 3. Views

### 3.1 Overview and Definition

### 3.2 Realization of the Generalization

### 3.3 Updates on Views

## 4. Window Functions

## 5. Sequences

## 6. (Other) Data Types and Built-In Functions

# Topics

## 1. Overview

## 2. Decomposing Complex Queries and Recursive Queries

## 3. Views

### 3.1 Overview and Definition

### 3.2 Realization of the Generalization

### 3.3 Updates on Views

## 4. Window Functions

## 5. Sequences

## 6. (Other) Data Types and Built-In Functions

# Views ...

- ... are **stored queries**, that are available as **virtual tables**
- ... originally used only for the **reading access**, but with certain constraints also for updates
- ... are constructed dynamically at each access.
- ... are not part of the physical schema

# Views ...

- use:
- concept for the adaptation to several users
  - simplification of queries
  - realization of the generalization

# Views ...

- use:
- concept for the adaptation to several users
  - simplification of queries
  - realization of the generalization

## DBMS proprietary:

- **materialized views** (ORACLE):  
pre-compiled non virtual query
- **indexed views** (SQL Server):  
additionally stored index

used to speed-up frequently executed queries



# Adjustment to Several Usergroups

## Example

due to data security in general only the exam list but not the grade are readable:

```
create view examineView as
select matrNr, lecNr, persNr
from examine;
```

# Simplification of Queries

## Example

avoiding constant use of join

*student* ⋈ *attend* ⋈ *lecture* ⋈ *professor*

```
create view
  StudProf(Sname, semester, title, Pname) as
select s.name, s.semester, v.title, p.name
from student s, attend h,
      lecture v, professor p
where s.matrNr=h.matrNr and
      h.lecNr=v.lecNr and
      v.heldBy = p.persNr;
```

# Simplification of Queries

## Example

avoiding constant use of join

*student* ⋈ *attend* ⋈ *lecture* ⋈ *professor*

```
create view
  StudProf(Sname, semester, title, Pname) as
select s.name, s.semester, v.title, p.name
from student s, attend h,
      lecture v, professor p
where s.matrNr=h.matrNr and
      h.lecNr=v.lecNr and
      v.heldBy = p.persNr;
```

find name and semester of all students of Sokrates

```
select distinct name, semester from StudProf
where Pname='Sokrates';
```

# Topics

## 1. Overview

## 2. Decomposing Complex Queries and Recursive Queries

## 3. Views

### 3.1 Overview and Definition

### 3.2 Realization of the Generalization

### 3.3 Updates on Views

## 4. Window Functions

## 5. Sequences

## 6. (Other) Data Types and Built-In Functions

# Realization of Generalization

realization of **inclusion** and **inheritance** via views:

- tuple of sub type should belong automatically to the super type
- attributes of the super type: should be inherited automatically

# Realization of Generalization

realization of **inclusion** and **inheritance** via views:

- tuple of sub type should belong automatically to the super type
- attributes of the super type: should be inherited automatically



define either sub type or super type as view  
decision based on access frequency

# Realization of Generalization

## Example (sub types as view)

relations for storing the information

(super type + additional information for sub types):

```
create table employee(  
    persNr integer not null,  
    name varchar(30) not null);  
create table profData (  
    persNr integer not null,  
    rank character(2),  
    room integer);  
create table assiData (  
    persNr integer not null,  
    field varchar(30),  
    boss integer);
```

# Realization of Generalization

## Example (sub types as view)

views with entire information about the sub types:

```
create view professor as select *  
from employee natural join profData;  
  
create view assistant as select *  
from employee natural join assiData;
```



# Realization of Generalization

## Example (sub types as view)

views with entire information about the sub types:

```
create view professor as select *  
from employee natural join profData;  
  
create view assistant as select *  
from employee natural join assiData;
```

only the relation employee as well as views professor and assistant visible

# Realization of Generalization

## Example (super types as view)

```
create table professor(  
    persNr integer not null,  
    name varchar(30) not null,  
    rank character(2),  
    room integer);  
  
create table assistant(  
    persNr integer not null,  
    name varchar(30) not null,  
    field varchar(30),  
    boss integer);  
  
create table otherEmpl(  
    persNr integer not null,  
    name varchar(30) not null);
```

# Realization of Generalization

## Example (super types as view)

```
create view employee as
(select persNr, name from professor) union
(select persNr, name from assistant) union
(select * from otherEmpl);
```

idea: only relations professor and assistant and the view employee visible

# Topics

## 1. Overview

## 2. Decomposing Complex Queries and Recursive Queries

## 3. Views

### 3.1 Overview and Definition

### 3.2 Realization of the Generalization

### 3.3 Updates on Views

## 4. Window Functions

## 5. Sequences

## 6. (Other) Data Types and Built-In Functions

# Updates on Views

problematic: `insert`, `update`, `delete` on a view

## Example (views incapable of update)

```
create view
  howDifficultAsExaminer(persNr, average) as
select persNr, avg(grade)
from examine
group by persNr;
```

# Updates on Views

problematic: `insert`, `update`, `delete` on a view

## Example (views incapable of update)

```
create view
  howDifficultAsExaminer(persNr, average) as
select persNr, avg(grade)
from examine
group by persNr;
```

```
create view lecturesView as
select title, SWS, name
from lectures, professor
where heldBy = persNr;
```

# Updates on Views

**restriction** of views capable of update in SQL to (SQL-92):

- only **one basic relation**
- **key** has to exist
- no aggregate functions, grouping and duplicate elimination

# Updates on Views

restriction of views capable of update in SQL to (SQL-92):

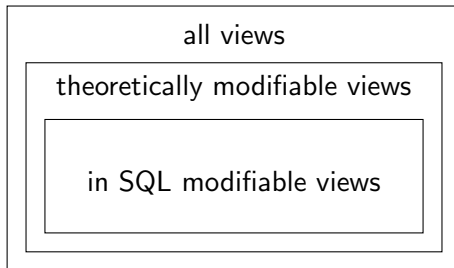
- only one basic relation
- key has to exist
- no aggregate functions, grouping and duplicate elimination

extension with SQL-99 (idea: consider key):

- a field of a view can be updated when
  - it can be assigned uniquely to a basic relation and
  - the corresponding key is part of the view
- distinction between views capable of update and views that allows for inserting new tuples



# Updates on Views



# Topics

1. Overview
2. Decomposing Complex Queries and Recursive Queries
3. Views
4. Window Functions
5. Sequences
6. (Other) Data Types and Built-In Functions
7. Learning Objectives

# From an Exercise...

## Example

```
SELECT floorNr, type
FROM   hive
WHERE  (SELECT COUNT(*) OVER ()
        FROM workerBee NATURAL JOIN pollinate
        WHERE worksinType = type AND ...
        GROUP BY fieldID, town LIMIT 1)
        < 10 AND
...;
```

# Window Functions: OVER ()

- allows for the calculation of (for instance aggregate-) functions over defined window
- grouping of rows not necessary

# Window Functions: OVER ( )

- allows for the calculation of (for instance aggregate-) functions over defined window
- grouping of rows not necessary

```
function_name ( args )  
  [ FILTER ( WHERE filter_clause ) ]  
  OVER ( window_definition )
```

- in SELECT or ORDER BY

here only basics, there are also other functionalities

# Aggregate Functions without Grouping

## Example

find for all students the number of attended lectures

```
SELECT DISTINCT matrnr,  
       count(*) OVER (PARTITION BY matrnr)  
       as number  
FROM attend  
ORDER BY number DESC;
```

# Avoiding Nested SELECT

## Example

find for all students all attended lectures and the number of attended lectures

```
SELECT matrnr, lecNr,  
       count(*) OVER (PARTITION BY matrnr)  
       as number  
FROM attend  
ORDER BY number DESC;
```

# Evaluation Order of Window Functions

- evaluation logically after WHERE, GROUP BY, HAVING
- corresponding removed rows are not considered



# Evaluation Order of Window Functions

- evaluation logically after WHERE, GROUP BY, HAVING
- corresponding removed rows are not considered

## Example

find for all students the number of attended lectures with lecNr bigger than 5001

```
SELECT DISTINCT matrnr,  
       count(*) OVER (PARTITION BY matrnr)  
       as number  
FROM attend  
WHERE lecNr > 5001  
ORDER BY number DESC;
```

# Special Window Functions

- in addition to “normal” functions as window functions there are also extra functions
- have to be used via `OVER`
- allow for the computation of values within a partition
- for instance `row_number()`, `rank()`, ...

## Example: rank()

- assigns a “ranking” to each entry

### Example

find for each semester those students that attend the most lectures

```
SELECT * FROM (  
  SELECT matrnr, count(*) as numberVO, sem,  
    rank() OVER  
      (PARTITION BY sem ORDER BY count(*) DESC)  
      as pos  
  FROM attend NATURAL JOIN student  
  GROUP BY matrnr, sem  
  ORDER BY numberVo DESC) as subq WHERE pos=1;
```

# Window Frame

- in addition to the entire partition there is a window frame for each row
- subset of partition, depending on sorting:
  - without ORDER BY: entire partition
  - with ORDER BY: all rows from the beginning until current row, additionally all rows with same ORDER BY value

## Example

```
SELECT matrnr, lecNr, sws,  
       sum(sws) OVER (PARTITION BY matrnr  
                      ORDER BY lecNr)  
FROM attend NATURAL JOIN lectures  
ORDER BY matrnr;
```

# Topics

1. Overview
2. Decomposing Complex Queries and Recursive Queries
3. Views
4. Window Functions
5. Sequences
6. (Other) Data Types and Built-In Functions
7. Learning Objectives

# Idea

## problem:

- in some cases **unique** (integer) values are needed, whose **specific value is not relevant**
- example: invoice numbers, visitor numbers, numbering of entries, ...

# Idea

## problem:

- in some cases **unique** (integer) values are needed, whose **specific value is not relevant**
- example: invoice numbers, visitor numbers, numbering of entries, ...

## solution (in SQL): **SEQUENCE**

- constructs **numerical sequence** based on given properties (domain, step size, ascending/descending)
- **no value** is assigned **twice**
- in some circumstances numbers **cannot** be assigned, or cannot be assigned **in chronological order**
- modifications are kept also with ROLLBACK

# Constructing and Using Sequences

```
CREATE SEQUENCE name  
[INCREMENT [ BY ] increment]  
[MINVALUE minvalue | NO MINVALUE]  
[MAXVALUE maxvalue | NO MAXVALUE]  
[START [WITH] start] [CACHE cache]  
[[NO] CYCLE]
```



# Constructing and Using Sequences

```
CREATE SEQUENCE name
[INCREMENT [ BY ] increment]
[MINVALUE minvalue | NO MINVALUE]
[MAXVALUE maxvalue | NO MAXVALUE]
[START [WITH] start] [CACHE cache]
[[NO] CYCLE]
```

## Zugriff:

`nextval(sname)` increases the sequence and returns current value

`currval(sname)` returns the last value which was read with  
`nextval` in the session

`setval(sname, val)` sets the counter to the given value

# SEQUENCE (Example)

## Example

```
CREATE SEQUENCE persnr_seq
  START WITH 2200 INCREMENT BY 5
  MINVALUE 2200 MAXVALUE 15000
  NO CYCLE;

INSERT INTO professor VALUES(
  nextval('persnr_seq'), 'Nietzsche','C3',205);
INSERT INTO professor VALUES(
  nextval('persnr_seq'), 'Sun Tzu','C3',206);
SELECT curval('persnr_seq');
```

# Topics

1. Overview
2. Decomposing Complex Queries and Recursive Queries
3. Views
4. Window Functions
5. Sequences
6. (Other) Data Types and Built-In Functions
7. Learning Objectives

# (Postgre)SQL Data Types

- numeric data types: `integer`, `int`, `numeric(p,s)`, ...
- strings: `char(n)`, `varchar(n)`, `text`, ...
- date and time data types: `date`, `timestamp`, ...
- enumerating types: `ENUM`
- geometric data types: `point`, `line`, `circle`, ...
- XML and JSON data types
- arrays
- ...

# (Postgre)SQL Data Types

- numeric data types: integer, int, numeric(p,s), ...
- strings: char(n), varchar(n), text, ...
- date and time data types: date, timestamp, ...
- enumerating types: ENUM
- geometric data types: point, line, circle, ...
- XML and JSON data types
- arrays
- ...

## Example (ENUM data type)

```
CREATE TYPE rank AS ENUM ( 'C2', 'C3', 'C4' );
```

# Functions

we have already learned about some built-in functions

- **aggregate functions**  
(compute a value over several rows)
- occasionally also **single row functions**

## **other single row functions:**

- character functions
- numeric functions
- converting functions
- date functions

# Character Functions

- concatenation: of strings, for instance: `SELECT 'Postgre' || 'SQL'`
- `lower()` and `upper()`:  
converting to lower- and upper-case characters
- `length()`:  
returns the length of a strings.
- `substring(string from for)`:  
`substring('Thomas' from 2 for 3)` – returns 'hom'
- `replace (string from to)`:  
`replace('abcdefcdg', 'cd', 'XX')` – returns 'abXXefXXg'

# Numeric Functions

- „everything that can be done with a calculator“, for instance
  - `sqrt(x)`, `power(x,y)`
  - `exp(x)`, `ln(x)`, `log(b, x)`
  - `cos(x)`, `sin(x)`, `tan(x)`, `atan(x)`, etc.
  - `abs(x)`, `sign(x)`
- `round()`:  
rounding (optional: to a certain number of decimal places)
  - `round(105.75)` – returns 106
  - `round(105.75, 1)` – returns 105.8
  - `round(105.75, -1)` – returns 110



# Converting Functions

- purpose:  
transforming values between different data types
- for instance: `to_char()`, `to_number()`
  - `to_char(17)` – returns '17'
  - `to_char(12345.678, '99,999.99')` – returns '12,345.68'
  - `to_number('17')` – returns 17
  - `to_number('-12,345.67', '99,999.99')` – returns -12345.67
  - `to_number('123.45', '99.99')` – returns error!
  - `to_number('123.49', '999.9')` – returns 123.4
- many other functions

# Data Type DATE

- used for storing **dates** (day, month, year)
- no time (SQL standard)
- **CURRENT\_DATE** returns current date  
(according to SQL standard: begin of a transaction)
- `to_char` and `to_date` offer several **formatting possibilities**

## to\_char() Function

cast of DATE to a string

### Example (to\_char())

```
SELECT to_char(birthDate) FROM professor;  
SELECT to_char(sent) FROM order;
```

## to\_char() Function

cast of DATE to a string

### Example (to\_char())

```
SELECT to_char(birthDate) FROM professor;  
SELECT to_char(sent) FROM order;
```

### Example (formatting via to\_char())

```
to_char(birthDate, 'MONTH DD, YYYY')  
to_char(birthDate, 'DD-MON-YYYY')  
to_char(birthDate, 'Day, DD.MM.YY')  
to_char(CURRENT_DATE, 'DD-MON-YYYY')
```

## TO\_DATE() Function

cast of a string to DATE.

### Example (to\_char())

```
UPDATE clients SET  
  born = TO_DATE('12-JUN-1976')  
WHERE clientNr = 1001;
```

## TO\_DATE() Function

cast of a string to DATE.

### Example (to\_char())

```
UPDATE clients SET
  born = TO_DATE('12-JUN-1976')
WHERE clientNr = 1001;
```

### Example (formatting via to\_char())

```
TO_DATE('12-JUN-1976')
TO_DATE('12.06.1976', 'DD.MM.YYYY')
TO_DATE('October 3, 1974', 'Month DD, YYYY')
```

## AGE() Function

difference between two dates given as years/months/days

### Example

```
SELECT AGE(CURRENT_DATE ,  
           TO_DATE('01 Sep 2012', 'DD Mon YYYY'));  
-> "5 years 2 mons 14 days"  
    (entered 15.11.2017)
```

```
SELECT AGE(  
           TO_DATE('01 Nov 2012', 'DD Mon YYYY'),  
           TO_DATE('05 Dec 2013', 'DD Mon YYYY'));  
-> "-1 years -1 mons -4 days"
```

# Data Type TIME and TIMESTAMP

**TIME** for storing the time

**TIMESTAMP** for storing date and time

**CURRENT\_TIME** resp. **CURRENT\_TIMESTAMP** returns current time  
resp. time and date  
(according to SQL standard: begin of transaction)

**TO\_CHAR** resp. **TO\_TIMESTAMP** functions offer several formatting possibilities



## EXTRACT() Function

enables reading parts (day, month, ...) of a date

### Example

```
SELECT EXTRACT (YEAR FROM Datum) AS year,  
       EXTRACT (MONTH FROM Datum) AS month,  
       EXTRACT (DAY FROM Datum) AS day  
FROM orders WHERE clientNr = 1003;
```

# Topics

1. Overview
2. Decomposing Complex Queries and Recursive Queries
3. Views
4. Window Functions
5. Sequences
6. (Other) Data Types and Built-In Functions
7. Learning Objectives

# Learning Objectives

- How can a (longer) SQL query be structured?
- How do recursive SQL SELECT statements work? (understanding and composing)
- What are views?
- What are views used for?
- How are they constructed?
- How do UPDATES behave on views?
- What is a SEQUENCE, and how do we construct and use it?